Question 1

Question Type: MultipleChoice

A company stores data in an Amazon S3 bucket.

Which task is the responsibility of AWS?

Options:

A- Configure an S3 Lifecycle policy.
B- Activate S3 Versioning.
C- Configure S3 bucket policies.
D- Protect the infrastructure that supports S3 storage.

Answer:

D

Explanation:

According to the AWS Shared Responsibility Model, AWS is responsible for protecting the infrastructure that runs all of the services offered in the AWS Cloud, including Amazon S3. This infrastructure includes hardware, software, networking, and facilities that run AWS services.

A . Configure an S3 Lifecycle policy: Incorrect, as configuring S3 Lifecycle policies to manage object lifecycle (e.g., transitioning objects to different storage classes or deleting them after a certain period) is the customer's responsibility.

B . Activate S3 Versioning: Incorrect, as enabling S3 Versioning is a customer responsibility for managing data protection.

C . Configure S3 bucket policies: Incorrect, as setting and managing S3 bucket policies to control access is the customer's responsibility.

AWS Cloud References:

AWS Shared Responsibility Model

Amazon S3

Question 2

Which Amazon S3 storage class is MOST cost-effective for unknown access patterns?

Options:

A- S3 Standard
B- S3 Standard-Infrequent Access (S3 Standard-IA)
C- S3 One Zone-Infrequent Access (S3 One Zone-IA)
D- S3 Intelligent-Tiering

Answer:

D

Explanation:

Understanding S3 Intelligent-Tiering: S3 Intelligent-Tiering is designed to optimize costs by automatically moving data to the most cost-effective access tier based on changing access patterns. It is ideal for data with unknown or unpredictable access patterns.

Why S3 Intelligent-Tiering is Cost-Effective:

Automatic Tiering: Moves data between two access tiers (frequent and infrequent access) based on changing access patterns, optimizing storage costs without performance impact.

No Retrieval Fees: Unlike other storage classes, there are no retrieval fees in Intelligent-Tiering, making it cost-effective for data with unpredictable access patterns.

Monitoring and Automation: Automatically monitors access patterns and transitions data, reducing the need for manual intervention.

When to Use S3 Intelligent-Tiering:

Unpredictable Access Patterns: Ideal for datasets where the access frequency cannot be determined or changes frequently.

Cost Optimization: For organizations looking to minimize storage costs without sacrificing performance or requiring manual intervention to move data between tiers.

References:

Amazon S3 Intelligent-Tiering

Amazon S3 Storage Classes

Question 3

Question Type: MultipleChoice

A developer needs to maintain a development environment infrastructure and a production environment infrastructure in a repeatable fashion Which AWS service should the developer use to meet these requirements?

Options:

A- AWS Ground Station

B- AWS Shield

C- AWS loT Device Defender

D- AWS CloudFormation

Answer:

D

Explanation:

AWS CloudFormation is a service that allows developers to model and provision their AWS infrastructure in a repeatable and declarative way, using code and templates. AWS CloudFormation enables developers to define the resources they need for their development and production environments, such as compute, storage, network, and application services, and automate their creation and configuration.AWS CloudFormation also provides features such as change sets, nested stacks, and rollback triggers to help developers manage and update their infrastructure safely and efficiently12.References:

AWS CloudFormation

What is AWS CloudFormation?

Question 4

Question Type: MultipleChoice

Which option is an environment that consists of one or more data centers?

Options:

A- Amazon CloudFront
B- Availability Zone
C- VPC
D- AWS Outposts

Answer:

B

Explanation:

Understanding Availability Zones (AZs): An Availability Zone is a distinct location within an AWS region that is engineered to be isolated from failures in other AZs.

Characteristics of Availability Zones:

Data Centers: Each AZ consists of one or more discrete data centers with redundant power, networking, and connectivity.

High Availability: AZs are designed for high availability, providing low-latency network connections to other zones in the same region.

Fault Isolation: They provide fault isolation and are used to deploy applications and services to ensure high availability and reliability.

Use Cases for Availability Zones:

Multi-AZ Deployments: For services like RDS, deploying in multiple AZs ensures fault tolerance.

Disaster Recovery: Setting up resources in multiple AZs helps in quick recovery from failures.

Load Balancing: Distributing traffic across AZs using Elastic Load Balancing ensures optimal performance and availability.

References:

AWS Global Infrastructure

Understanding AWS Regions and Availability Zones

Question 5

Question Type: MultipleChoice

Which AWS Cloud design principle is a company using when the company implements AWS CloudTrail?

Options:

A- Activate traceability.

B- Use serverless compute architectures.

C- Perform operations as code.

D- Go global in minutes.

Answer:

A

Explanation:

By implementing AWS CloudTrail, a company is adhering to the AWS Cloud design principle of activating traceability. AWS CloudTrail provides detailed logs of all API calls made in an AWS account, which helps monitor, troubleshoot, and detect unusual activity, thereby improving security and compliance. This supports the principle of 'activating traceability' by enabling continuous monitoring and auditing of all actions and changes within the AWS environment.

B . Use serverless compute architectures: Incorrect, as this principle encourages the use of managed services that handle infrastructure, such as AWS Lambda, and is not directly related to CloudTrail.

C . Perform operations as code: Incorrect, as this principle emphasizes the use of code and automation for infrastructure management.

D . Go global in minutes: Incorrect, as this principle relates to the global deployment of applications and services.

AWS Cloud References:

AWS Well-Architected Framework

AWS CloudTrail

Question 6

MultipleChoice

Which AWS Cloud benefit is shown by an architecture's ability to withstand failures with minimal downtime?

Options:

A- Agility
B- Elasticity
C- Scalability
D- High availability

Answer:

D

Explanation:

Understanding High Availability: High availability (HA) refers to systems that are durable and likely to operate continuously without failure for a long time. HA ensures that an architecture can withstand failures with minimal downtime.

Importance of High Availability:

Redundancy: Systems are designed with redundancy to prevent single points of failure.

Fault Tolerance: Ensures that failures do not result in significant downtime, maintaining service continuity.

Automated Recovery: Utilizes automated recovery mechanisms to quickly restore services in the event of a failure.

AWS Services for High Availability:

Multi-AZ Deployments: Services like RDS, DynamoDB, and others support Multi-AZ deployments for fault tolerance.

Elastic Load Balancing: Distributes traffic across multiple instances or availability zones to ensure no single point of failure.

Auto Scaling: Automatically adjusts the number of instances based on demand, ensuring availability even during traffic spikes.

References:

AWS Well-Architected Framework: Reliability

Question 7

Question Type: MultipleChoice

Which AWS service or tool helps users visualize, understand, and manage spending and usage over time?

Options:

A- AWS Organizations

B- AWS Pricing Calculator

C- AWS Cost Explorer

D- AWS Service Catalog

Answer:

C

Explanation:

AWS Cost Explorer is the AWS service or tool that helps users visualize, understand, and manage spending and usage over time. AWS Cost Explorer is a web-based interface that allows users to access interactive graphs and tables that display their AWS costs and usage data. Users can create custom reports that analyze cost and usage data by various dimensions, such as service, region, account, tag, and more. Users can also view historical data for up to the last 12 months, forecast future costs for up to the next 12 months, and get recommendations for cost optimization. AWS Cost Explorer also provides preconfigured views that show common cost and usage scenarios, such as monthly spend by service, daily spend by linked account, and Reserved Instance utilization. Users can use AWS Cost Explorer to monitor their AWS spending and usage trends, identify cost drivers and anomalies, and optimize their resource allocation and budget planning.References:Cloud Cost Analysis - AWS Cost Explorer - AWS,Analyzing your costs with AWS Cost Explorer

To Get Premium Files for CLF-C02 Visit

https://www.p2pexams.com/products/clf-c02

For More Free Questions Visit

https://www.p2pexams.com/amazon/pdf/clf-c02